

# Missing optimizations in node-based containers

ACCU 2017  
April 28, 2017

Elliot Goodrich, Software Engineer  
[egoodrich4@bloomberg.net](mailto:egoodrich4@bloomberg.net)

techatbloomberg.com


© 2017 Bloomberg Finance L.P. All rights reserved.

Bloomberg

**How fast can we reverse a doubly linked list?**


# std::list::reverse

```
struct node {  
    node* next;  
    node* prev;  
};  
  
void reverse(node* head) {  
    node* curr = head;  
  
    do {  
        std::swap(curr->prev, curr->next);  
        curr = curr->prev;  
    } while (curr != head);  
}
```




# std::list::reverse benchmark results

**Sequentially Placed Nodes** 119.46 million nodes per second



- Memory layout matters!
- Accessing memory from RAM is slow (100+ CPU cycles)
- Predictable memory access allows cache prefetcher to fetch memory in advance
- Nodes have been padded to 64 bytes
- i5 3750K CPU (2012)


**Randomly Placed Nodes** 10.47 million nodes per second



# Instruction-Level Parallelism

- Modern CPUs have **Out-of-Order Execution**
- Instructions can be executed when their inputs are ready, not based on the instruction order
- Independent instructions can be executed in parallel (still on one thread)

```
int f(int* x, int y) {  
    int a = *x + 2;  
    int b = y - 1;  
    int c = y + 1;  
    return a * b * c;  
}
```



If  $*x$  is not in the cache, this operation dominates



**Goal: Remove data dependencies!**


# Data dependencies in `std::list::reverse`

```
void reverse(node* head) {  
    node* curr = head;  
  
    do {  
        std::swap(curr->prev, curr->next);  
        curr = curr->prev;  
    } while (curr != head);  
}
```

- Pointer chasing has data dependencies
- Modern CPUs have **Memory-Level Parallelism**
- Idea – follow the prev and next pointers together

# std::list::reverse (unrolled)

```
void reverse(node* head) {  
    node* fwd = head;  
    node* bck = head->prev;  
    if (fwd == bck) return;  
    while (true) {  
        node* fwd2 = fwd->next;  
        node* bck2 = bck->prev;  
        std::swap(fwd->next, fwd->prev);  
        std::swap(bck->next, bck->prev);  
  
        if (fwd2 == bck2) {  
            std::swap(fwd2->next, fwd2->prev);  
            return;  
        }  
        if (fwd2 == bck) return;  
  
        fwd = fwd2;  
        bck = bck2;  
    }  
}
```





# std::list::reverse (unrolled)

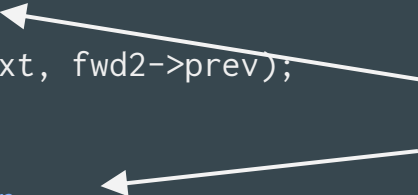
```
void reverse(node* head) {
    node* fwd = head;
    node* bck = head->prev;
    if (fwd == bck) return;
    while (true) {
        node* fwd2 = fwd->next;
        node* bck2 = bck->prev;
        std::swap(fwd->next, fwd->prev);
        std::swap(bck->next, bck->prev);

        if (fwd2 == bck2) {
            std::swap(fwd2->next, fwd2->prev);
            return;
        }
        if (fwd2 == bck) return;

        fwd = fwd2;
        bck = bck2;
    }
}
```

- Twice as much work per iteration
- CPU can fetch fwd2 and bck2 in parallel since they are independent instructions
- **The optimisation is from the CPU!**
- Language agnostic

If we know the size of the linked list (required in C++11 onwards) then one of these branches can be removed



# std::list::reverse (unrolled single if)

```
void reverse(node* head, int size) {
    node* fwd = head;
    node* bck = head->prev;
    if (size % 2 == 1) {
        std::swap(fwd->prev, fwd->next);
        fwd = fwd->prev;
    }
    if (size == 1) return;

    while (true) {
        node* fwd2 = fwd->next;
        node* bck2 = bck->prev;
        std::swap(fwd->next, fwd->prev);
        std::swap(bck->next, bck->prev);

        if (fwd2 == bck) return;
        fwd = fwd2;
        bck = bck2;
    }
}
```

# Benchmarks

	Sequential	Speedup	Random	Speedup
reverse	119.46 M	1.00x	10.47 M	1.00x
unrolled	125.42 M	1.05x	19.23 M	1.84x
unrolled single if	125.42 M	1.05x	19.23 M	1.84x

# Other functions that could equally benefit

- `list::merge`
- `list::splice` ( $O(N)$  overload splicing one iterator range to another list)
- `list::remove`/`list::remove_if`
- `list::merge`/`map::merge`
- `list::unique`
- `list::sort`
- `distance` (on `list`/`map` iterators)
- `any_of`/`all_of`/`none_of` (on `list`/`map` iterators)
- Any hash table (e.g. `unordered_map`) implementation that uses a doubly-linked list to solve collisions

# Why does it fall short of 2x improvement?

```
node* fwd2 = fwd->next;
```

```
node* bck2 = bck->prev;
```

```
std::swap(fwd->next, fwd->prev);
```

```
std::swap(bck->next, bck->prev);
```



Read-After-Write (RAW) dependency  
between these swaps

- fwd and bck could alias
- CPU must wait for first swap before executing the second

# Optimizing swaps

read

write

```
node* fwd2 = fwd->next;
node* bck2 = bck->prev;
// std::swap(fwd->next, fwd->prev);
fwd->next = fwd->prev;
fwd->prev = fwd2;
// std::swap(bck->next, bck->prev);
bck->prev = bck->next;
bck->next = bck2;
```

asm

mov  
mov  
mov  
mov  
mov  
mov  
mov

```
rax, QWORD PTR [rdi+8]
rdx, QWORD PTR [rdi]
QWORD PTR [rdi], rax
QWORD PTR [rdi+8], rdx
rdx, QWORD PTR [rsi]
rax, QWORD PTR [rsi+8]
QWORD PTR [rsi+8], rdx
QWORD PTR [rsi], rax
```

rewrite

```
node* fwd2 = fwd->next;
node* bck2 = bck->prev;
node* tmp  = bck->next;
fwd->next = fwd->prev;
fwd->prev = fwd2;
bck->next = bck2;
bck->prev = tmp;
```

asm

mov  
mov  
mov  
mov  
mov  
mov  
mov  
mov

```
rcx, QWORD PTR [rdi+8]
rdx, QWORD PTR [rsi]
rax, QWORD PTR [rsi+8]
r8, QWORD PTR [rdi]
QWORD PTR [rdi], rcx
QWORD PTR [rdi+8], r8
QWORD PTR [rsi+8], rdx
QWORD PTR [rsi], rax
```

# std::list::reverse (double\_swap)

```
void reverse(node* head, int size) {
    node* fwd = head;
    node* bck = head->prev;
    if (size % 2 == 1) {
        std::swap(fwd->prev, fwd->next);
        fwd = fwd->prev;
    }
    if (size == 1) return;

    while (true) {
        node* fwd2 = fwd->next;
        node* bck2 = bck->prev;
        → double_swap(fwd, bck);

        if (fwd2 == bck) return;
        fwd = fwd2;
        bck = bck2;
    }
}
```

# Benchmarks

	Sequential	Speedup	Random	Speedup
reverse	119.46 M	1.00x	10.47 M	1.00x
unrolled	125.42 M	1.05x	19.23 M	1.84x
unrolled single if	125.42 M	1.05x	19.23 M	1.84x
double_swap	128.21 M	1.07x	20.40 M	1.95x



# Other examples of unnecessary RAW dependencies

// Visual Studio

```
node* _Unlinknode(const_iterator _Where) {  
    node* _N = _Where.node();  
    _N->prev->next = _N->next;  
    _N->next->prev = _N->prev;  
    --this->size;  
    return _N;  
}
```

// Visual Studio

```
node* _Unlinknode(const_iterator _Where) {  
    node* _N = _Where.node();  
    node* _Next = _N->next;  
    node* _Prev = _N->prev;  
    _Prev->next = _Next;  
    _Next->prev = _Prev;  
    --this->size;  
    return _N;  
}
```

// libc++

```
void __unlink_nodes(node* __f, node* __l) {  
    __f->prev->next = __l->next;  
    __l->next->prev = __f->prev;  
}
```

// libc++

```
void __unlink_nodes(node* __f, node* __l) {  
    node* __next = __l->next;  
    node* __prev = __f->prev;  
    __prev->next = __next;  
    __next->prev = __prev;  
}
```

# Does the `restrict` keyword help?

- `restrict` qualifier (and `__restrict`/`__restrict__`) restricts **pointer aliasing**
- `restrict` would give more information to the compiler (e.g. our writes are not dependent)
- Success differs between compilers and what function you're optimizing
- Rewriting functions gives more consistent results

# Further (and slightly crazier) Optimizations

- Need more pointers to chase in parallel - requires making changes to the linked list
- If our linked list held a pointer to a node midway then this would give us 4 independent pointers to chase
  - Following next from head
  - Following prev from head
  - Following next from midway
  - Following prev from midway
- Problem! Can't update midway while splicing
- Require  $O(N)$  time to update midway, but splice within a list must be  $O(1)$

# Loopidly linked lists


- next pointer is unchanged
- prev pointer points to the node **two positions back** in the list

## Pros

- Independently chase two pointers backwards and one forward

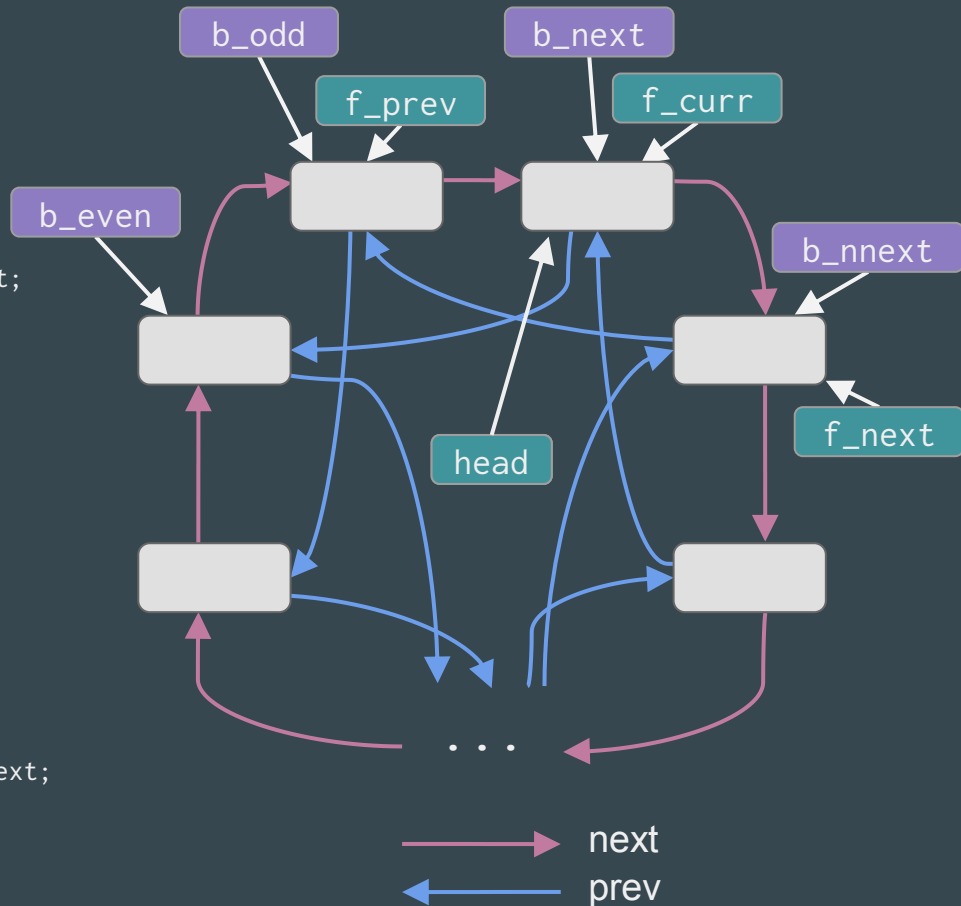
## Cons

- Given a node  $n$ , the node before it will be  $n \rightarrow \text{prev} \rightarrow \text{next}$ , which will take twice as long to access
- Have to modify two extra nodes when inserting or deleting



# Loopidly linked lists

```
void reverse(node* head, int size) {  
    node* f_prev = head->prev->next;  
    node* f_curr = head;  
    node* f_next = head->next;  
    node* b_nnext = f_next;    node* b_next = head;  
    node* b_even = head->prev; node* b_odd = b_even->next;  
    switch (size / 3) { /* process 0-2 items */  
  
        while (true) {  
            node* f_nnext = f_next->next;  
            node* b_prev = b_odd->prev;  
            node* b_pprev = b_even->prev;  
            } Independent loads  
  
            f_curr->next = f_prev; f_curr->prev = f_nnext;  
            b_odd->next = b_even; b_odd->prev = b_nnext;  
            b_even->next = b_prev; b_even->prev = b_next;  
  
            if (b_prev == f_curr) return;  
  
            f_prev = f_curr; f_curr = f_next; f_next = f_nnext;  
            b_nnext = b_odd; b_next = b_even;  
            b_odd = b_prev; b_even = b_pprev;  
        }  
    }
```



# Benchmarks

	Sequential	Speedup	Random	Speedup
reverse	119.46 M	1.00x	10.47 M	1.00x
unrolled	125.42 M	1.05x	19.23 M	1.84x
unrolled single if	125.42 M	1.05x	19.23 M	1.84x
double_swap	128.21 M	1.07x	20.40 M	1.95x
loopidly	128.21 M	1.07x	26.83 M	2.56x

# Loopidly linked lists

- `prev` pointer could skip three nodes (or four, or five...) giving more independent loads
- Or the `next` pointer could skip ahead as well
- Diminishing returns on skipping more nodes

# Thanks for listening!



Elliot Goodrich  
egoodrich4@bloomberg.net  
elliotgoodrich@gmail.com